

EXERCISES ON HIGH PERFORMANCE COMPUTING

Block 1: Memory Hierarchy

S. Hunold, J. L. Träff, and F. Versaci

*Parallel Computing Group
TU Wien*

15 April, 2013

Generalities

Instructions. There will be four (or maybe five) blocks of exercises and you are required to solve *two blocks* out of them.

- To solve this block, solve *one of the two* following exercises (preferably in C/C++).
- To avoid penalizations, observe *exactly* the provided guidelines.
- The implementation will be judged based on correctness and performance.
- It is *strictly forbidden* to copy/take inspiration from code from other people or taken from internet. You are required to *write your own code*.
- You are required to hand in (to the contact email) your source code (including the Makefile, excluding any input file you might have used for testing purposes) in a tar.gz or zip file.
- You are also required to attach a document (5-15 pages, containing no source code, to be provided in pdf format) which should contain roughly the following sections:
 1. Broad description of the implemented algorithm (1-3 pages)
 2. Some implementation details (1-3 pages)
 3. Presentation of experimental results (2-6 pages)
 4. Analysis of the results (2-4 pages)

Contact. If you need some clarifications for this block, contact `versaci@par.tuwien.ac.at`.

Deadline. 24 June 2013, 24:00.

Exercises

1. We are interested in computing the number of misses incurred by the LRU eviction policy, given an input address trace and for any possible buffer size.
 - The address trace is made of 64-bit little-endian unsigned integers, stored in a file (create your own files for testing purposes).
 - The name of the input file should be passed as an argument in the command line.
 - The program should be named `lru-misses` and should either come with a Makefile or bash script for compiling it.

- The program should provide as output in the stdout the number of misses for buffers of sizes 2^k , with $k \in \{0, 1, \dots, 64\}$.
- The number of misses should be printed in plain text, one per line.
- The expected behavior is thus the following:

```
$ make
$ ./lru-misses inputfile.bin
46842
23651
18454
...
```

- The algorithm you are required to implement is the one appearing in [1]
2. We want to sort an array of little-endian doubles, stored in a file, and save the sorted output into another file (both filenames to be passed as arguments in the command line).
 - The file might not fit into main memory (external sort) and should be mapped to virtual memory using `mmap` calls.
 - The program should be named `co-sort` and should either come with a Makefile or bash script for compiling it.
 - The expected behavior is thus the following:


```
$ make
$ ./co-sort inputfile.bin outputfile.bin
```
 - You are required to implement a cache oblivious sorting algorithm, choosing among the following ones
 - (a) Distribution sorting [5]
 - (b) Funnelsort [5]
 - (c) Lazy funnelsort [2,3]
 - (d) Proximity mergesort [4]
 - You are also required to compare the performance against a standard quick sort implementation (you can use C `qsort` from `stdlib` or C++ STL `sort`)

References

- [1] ALMÁSI, G., CAŞCAVAL, C., AND PADUA, D. A. Calculating stack distances efficiently. *SIGPLAN Not.* 38, 2 supplement (June 2002), 37–43.
- [2] BRODAL, G. S., AND FAGERBERG, R. Cache oblivious distribution sweeping. In *ICALP* (2002), pp. 426–438.
- [3] BRODAL, G. S., FAGERBERG, R., AND VINTHER, K. Engineering a cache-oblivious sorting algorithm. *ACM Journal of Experimental Algorithmics* 12 (2007).
- [4] FRANCESCHINI, G. Proximity mergesort: optimal in-place sorting in the cache-oblivious model. In *SODA* (2004), pp. 291–299.
- [5] FRIGO, M., LEISERSON, C. E., PROKOP, H., AND RAMACHANDRAN, S. Cache-oblivious algorithms. *ACM Trans. Algorithms* 8, 1 (Jan. 2012), 4:1–4:22.